

# Data Science Lecture 02

24.04.2015

Dr. Kashif Rasul

 kashif  @krasul  #167

Welcome to Data Science Lecture 2.

# Last Time

- Visualisation
- Tufte's Principles
- Choice of visualisations for 4 different data types
- Homework 0 🐏🚗🚪🐏

2

Last time we gave you an overview of information visualisation and did not really get into story telling.




3

And if we get time near the end of the course, we'll dive into story telling a bit more, but if you want to understand story telling at least from the point of view of comics, then all of Scott McCloud's books are worth reading. If you get one of his books then we recommend "Understanding Comics" from 1993.


The pedagogical potential of the comic form is vast and mostly untapped. This book explores the formal aspects of comics, the historical development of the medium, its fundamental vocabulary, and various ways in which these elements have been used. It expounds theoretical ideas about comics as an art form and medium of communication and story telling.

### What kind(s) of computer(s) do you own?



Desktop	8	11.3%
Laptop	63	88.7%
None	0	0%


### What operating system(s) do you run on your computer(s)?



OS X	21	29.6%
Linux	45	63.4%
Windows	29	40.8%
*BSD	0	0%
Other	0	0%


Here are the results of the initial student survey. Thank you for that. We will take into account the great feedback.

**What operating system(s) do you run on your mobile(s)?**



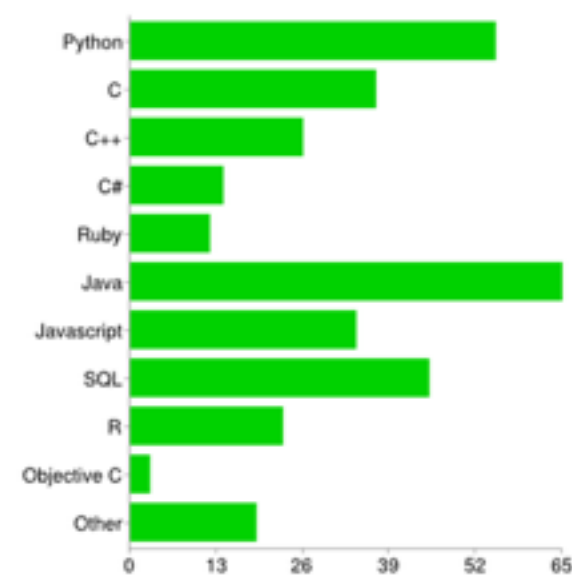
Android seems the dominant mobile OS here.

### How long have you been programming?



Less than 6 months	2	2.8%
Between 6 and 12 months	5	7%
1 to 3 years	12	16.9%
more than 3 years	52	73.2%


### What languages do you know?



Python	55	77.5%
C	37	52.1%
C++	26	36.6%
C#	14	19.7%
Ruby	12	16.9%
Java	65	91.5%
Javascript	34	47.9%
SQL	45	63.4%
R	23	32.4%
Objective C	3	4.2%
Other	19	26.8%

Ok so Python seems like a good choice for this course. No surprise here at that Java is the most popular language here.

**Overall, how comfortable are you with programming?**



1	2	2.8%
2	4	5.6%
3	17	23.9%
4	27	38%
5	21	29.6%

# Today

- Data Scrapping, Munging and Cleaning
- Live coding
- Exercise 1
- Homework 1



# Getting WWW data

- Typical workflow: Direct Download (check wiki)
- Make an API call (twitter, facebook etc.)
- Web scrapping: Homework 1

9

Let us turn our attention now to the problem of getting data from the WWW and cover some basics so we can be all on the same page.

When a website offers data then it is relatively easy to get the data in some standard format. On the github wiki there is a Data sources page with links to sources of data.

More and more often a service provides an API to get data.

Lastly, if you can't get the data in any other way, you might be forced to scrape the data.

```

{
  "offset": "0",
  "results": [
    {
      "body": "IF there is one thing that the past two weeks have shown, it's that America is about to enter a new age of talk. Even if there weren't going to be an election next November that would usher in a new administration, the United States, under President Bush, would be entering a new age of talk. Seven years of President Bush's Don't-Talk-to-Evil policy",
      "byline": "By HELENE COOPER",
      "date": "20071216",
      "title": "THE NATION; Look Who Talks To The Enemy",
      "url": "http:\\\\query.nytimes.com\\gst\\fullpage.html?res=9C04E4D81F31F935A25751C1A9619C8B63"
    },
    {
      "body": "JUST 24 hours after Hillary Clinton mowed down a skeptical Katie Couric with her certitude that she would win the Democratic nomination -- 'It will be me!' -- her husband showed exactly how she could lose it. By telling an Iowa audience on Tuesday night that he had opposed the Iraq war 'from the beginning,' Bill Clinton committed a double",
      "byline": "By FRANK RICH",
      "date": "20071202",
      "title": "OP-ED COLUMNIST: Who's Afraid of Barack Obama?"
    }
  ]
}

```

Typically an API call will return data in a format called JSON (JavaScript Object Notation). It corresponds to a mix between Dictionaries (key, values) and Lists.

It's popularity is growing and is the interface we will try to use most often.

Here is an example of the result of an API request to the New York Times.

```
http://api.nytimes.com/svc/search/v1/article?  
format=json&query=obama+syria&api-key=####
```

11

Now this particular JSON was received after we made a call to this Request URL.

Now this Web query string is made up of the URL and a search part which is separated by a '?' character. The search part is made up of a series of key value fields separated by the '&' character.

One final thing to note is that most API's require you to register with an API key. By changing the search part we can request terms we want on the fly.

A good API will also come with some documentation on the query terms and the type of responses one should accept. For example if I want to search for some term in the title of articles, I can look up the NewYork Times' API documentation.

# HTTP

- GET: used to retrieve data and have no other effect
- POST: server accept the item enclosed (forms)
- PUT: enclosed entity be stored (update name)
- DELETE: delete the specified resource

12

We mentioned a GET request, but to step back a bit, we need to talk about HTTP. HTTP is a application protocol for distributed, collaborative, hypermedia information systems. Web servers and web browsers communicate with each other via this protocol.

HTTP defines methods (sometimes referred to as verbs) to indicate the desired action to be performed on the identified resource. The four most common verbs are:

```
<?xml version="1.0" encoding="UTF-8"?>
<result_set>
  <status>
    OK
  </status>
  <copyright>
    Copyright (c) 2014 The New York Times Company.
  </copyright>
  <num_results>
    25
  </num_results>
  <last_modified>
    2014-04-11T13:07:16-04:00
  </last_modified>
  <results>
    <book>
      <list_name>
        Hardcover Fiction
      </list_name>
      <display_name>
        Hardcover Fiction
      </display_name>
      <updated>
        WEEKLY
      </updated>
    </book>
  </results>
</result_set>
```

Sometimes the response format is in XML and is characterised by opening and closing tags in a hierarchical or nested manner.

This is different from Json, but XML is extremely powerful in its own sense and more verbose and a bit more cumbersome to use for humans.



14

Finally we can always use Chrome's developer tools and look at the HTML source of any webpage which is online. Go to View -> Developer Tools.

HTML is a simplified version of XML with predefined tags. So for web scraping it is important that you understand HTML.

Your code will need to get web page via a GET request.

```

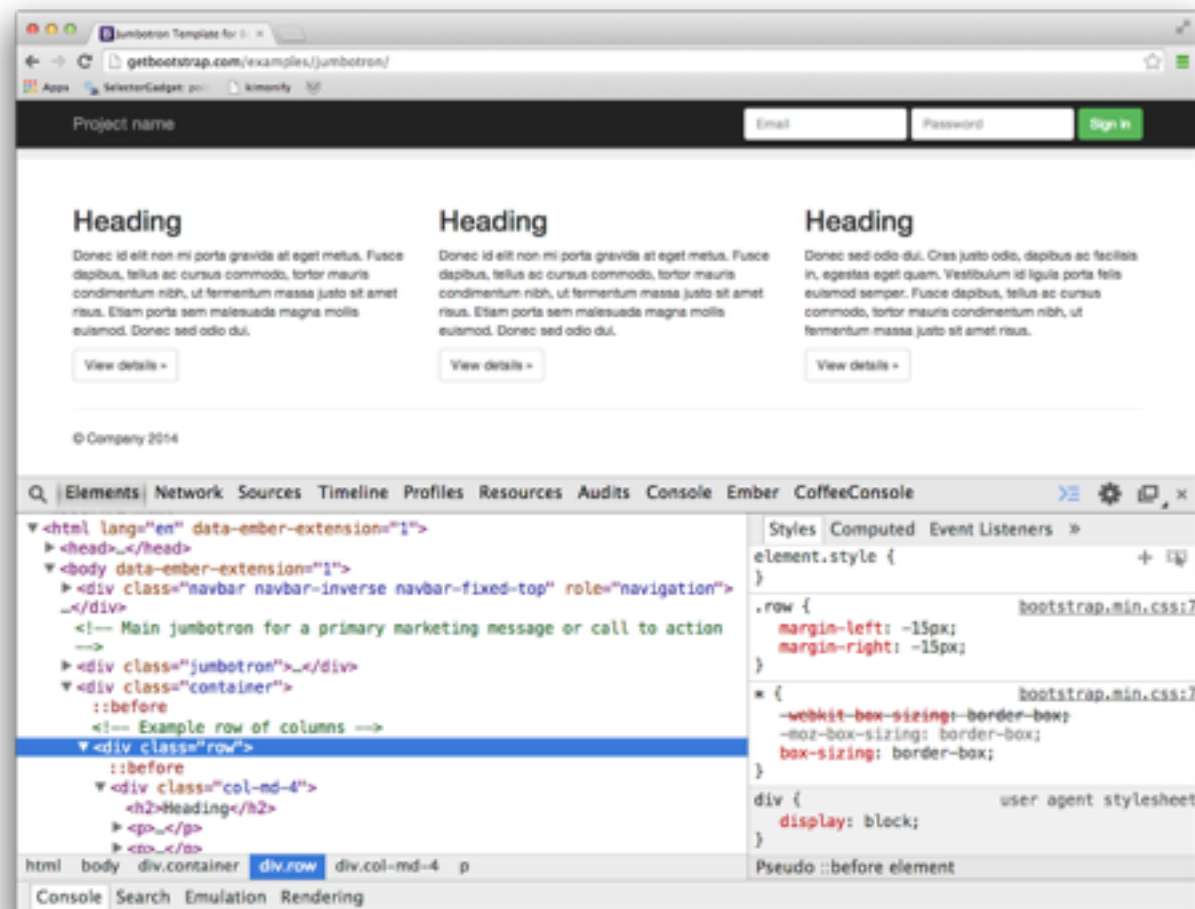
...
<div class="row">
  <div class="col-md-4">
    <h2>Heading</h2>
    <p>Donec id elit ui </p>
    <p><a class="btn" href="#" role="button">View details</a></p>
  </div>
  <div class="col-md-4">
    <h2>Heading</h2>
    <p>Donec id.</p>
    <p><a class="btn" href="#" role="button">View details</a></p>
  </div>
  <div class="col-md-4">
    <h2>Heading</h2>
    <p>Donec sed odio dui.</p>
    <p><a class="btn" href="#" role="button">View details</a></p>
  </div>
</div>
...

```

15

In looking at HTML source we have some predefined tags like 'div' together with class attributes like "row" etc. So the purpose of the class attribute is for styling via CSS and when a number of objects are alike for example the styling of a row or column we use the class attribute.

There is also the 'id' attribute which is used for styling specific elements of the page like the for example if a particular item in a list is non-editable, we might make it grey depending on it's 'id'.



16

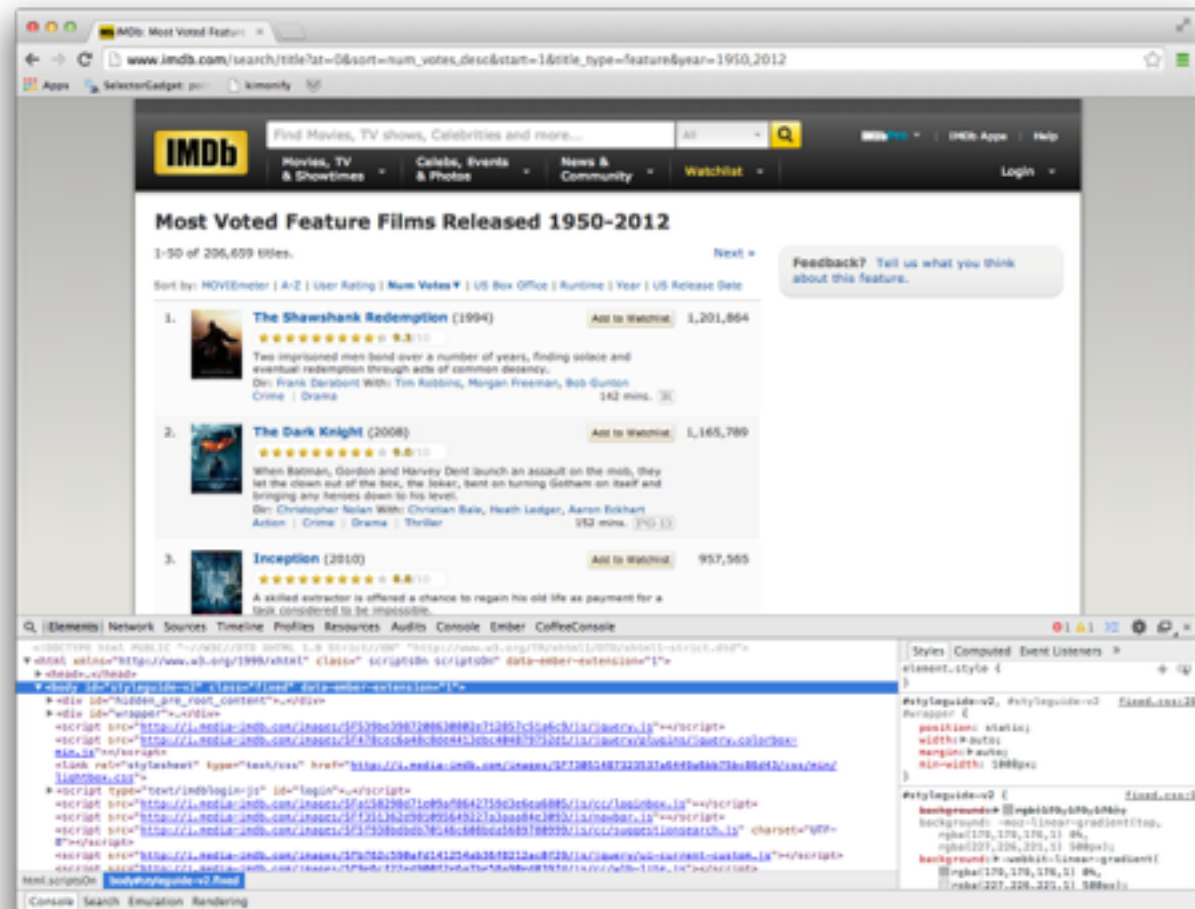
And this hierarchical or tree representation of the webpage is referred to as the DOM (Document Object Model). And the chrome inspector makes it easy to explore the DOM.

For scrapping we need to be able to parse the DOM, traverse it and extract appropriate elements from it that we need from this DOM, all from python.



# Ethics of Scrapping

- Scrape public websites only
- Non-sensitive information only
- If you do scrape try to anonymise the data
- Always cite your sources (url, date, time etc.)



18

So let's try to find and print the movie title, list of genres, runtime, and score of all movies on this page: "The most voted feature films released from 1950-2012" on IMDb.

```
import requests
from pattern import web
from bs4 import BeautifulSoup
```

In python scrapping is usually done with requests, pattern or BeautifulSoup. Let's import them all.

```
url = 'http://www.imdb.com/search/title?
sort=num_votes,desc&start=1&title_type=feature&year=1950,2012'

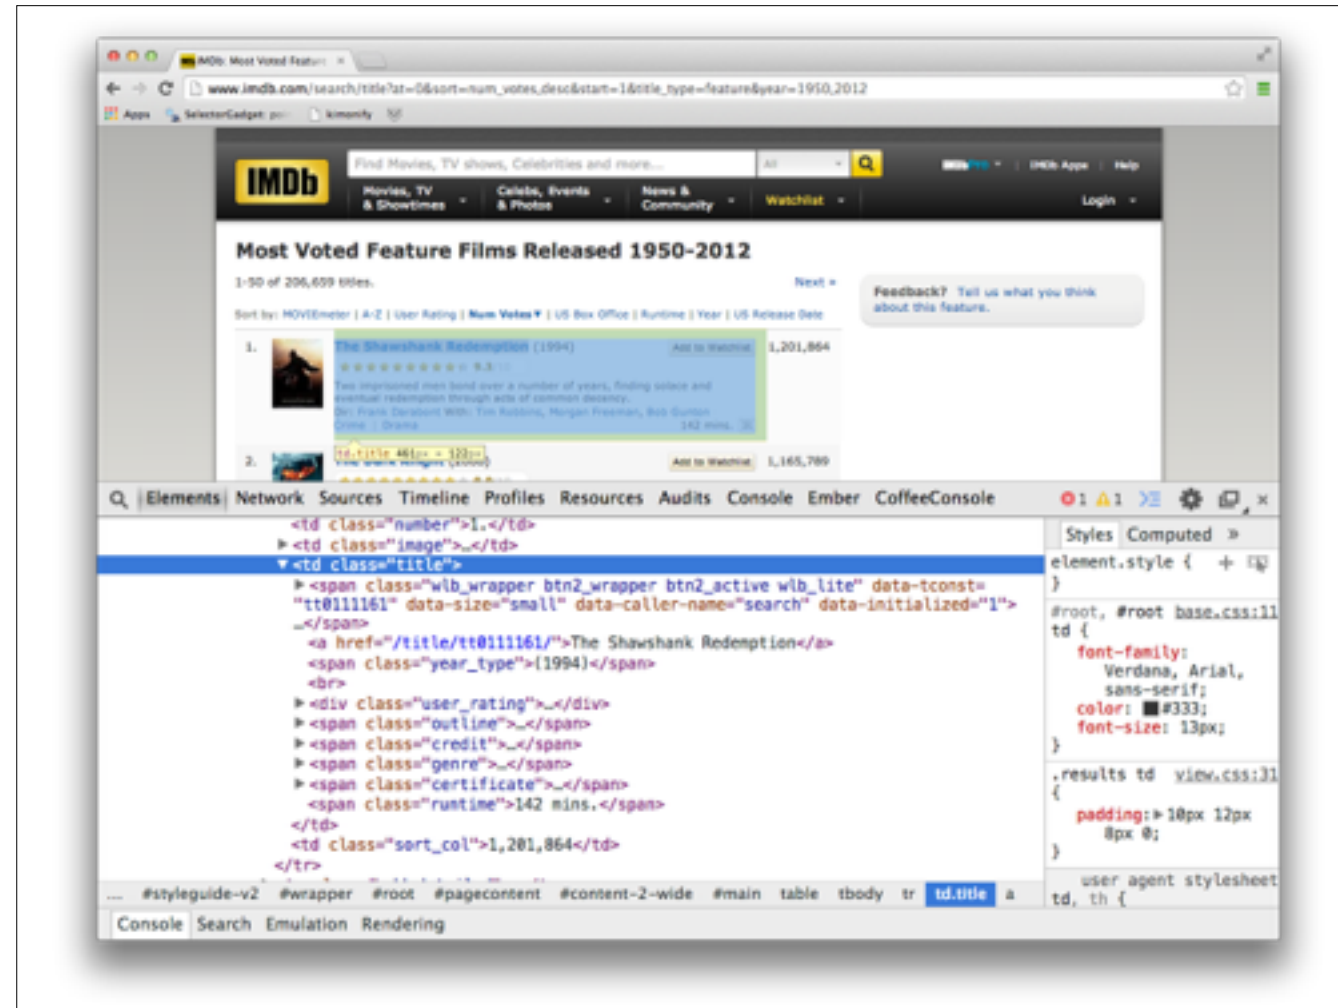
r = requests.get(url)
print r.url

http://www.imdb.com/search/title?
sort=num_votes,desc&start=1&title_type=feature&year=1950,2012

url = 'http://www.imdb.com/search/title'
params = dict(sort='num_votes,desc', start=1,
              title_type='feature', year='1950,2012')

r = requests.get(url, params=params)
print r.url # notice it constructs the full url for you
```

Well we can explicitly to a GET request on this url. Or we can construct the url via a base URL together with a dictionary of params.



Each movie is in the DOM element table row 'tr'->'td.title'. Inside each 'td.title' we have the title in the link tag, the genres are in the span class 'genre' etc. So we need to loop over all the 'td.title' tags and get the information we need.

```

dom = web.Element(r.text)

for movie in dom.by_tag('td.title'):
    title = movie.by_tag('a')[0].content
    genres = movie.by_tag('span.genre')[0].by_tag('a')
    genres = [g.content for g in genres]
    runtime = movie.by_tag('span.runtime')[0].content
    rating = movie.by_tag('span.value')[0].content
    print title, genres, runtime, rating

The Shawshank Redemption [u'Crime', u'Drama'] 142 mins. 9.3
The Dark Knight [u'Action', u'Crime', u'Drama', u'Thriller'] 152
mins. 9.0
Inception [u'Action', u'Adventure', u'Mystery', u'Sci-Fi',
u'Thriller'] 148 mins. 8.8
Pulp Fiction [u'Crime', u'Drama', u'Thriller'] 154 mins. 9.0
Fight Club [u'Drama'] 139 mins. 8.9
...

```

So with patterns we can get the particular DOM by it's css tag. Some tags like the link has the information inside it's content and so we use '.content' to get that information. Most of the work of scrapping is iterating this process for the tag you want.

Note the list of genres are created via the python's syntactic sugar which is equivalent to doing:

```

genres = []
for g in genres
    genres.append(g.content)

```

```

bs = BeautifulSoup(r.text)

for movie in bs.findAll('td', 'title'):
    title = movie.find('a').contents[0]
    genres = movie.find('span', 'genre').findAll('a')
    genres = [g.contents[0] for g in genres]
    runtime = movie.find('span', 'runtime').contents[0]
    rating = movie.find('span', 'value').contents[0]
    print title, genres, runtime, rating

The Shawshank Redemption [u'Crime', u'Drama'] 142 mins. 9.3
The Dark Knight [u'Action', u'Crime', u'Drama', u'Thriller'] 152
mins. 9.0
Inception [u'Action', u'Adventure', u'Mystery', u'Sci-Fi',
u'Thriller'] 148 mins. 8
...

```

In BeautifulSoup we use the findAll() and find() functions to get the DOM elements we want.

# EDA: Workflow

- Build a data frame
- Clean the data frame
  - each row: describes a single object
  - each column: a property which is numeric and atomic
- Global properties
- Group properties

24

The basic workflow is to build a data frame, clean up the data frame, explore the global properties and the group properties of the data.

This process transforms your data into a format which is easier to work with, gives you a basic overview of the data's properties, and likely generates several questions for you to follow up in subsequent analysis.





Cool that was easy. Well this is the feeling one gets when doing Data wrangling and exploration.

```
!head data/imdb_top_10000.txt

tt0111161 The Shawshank Redemption (1994) 1994 9.2 619479 142
mins. Crime|Drama
tt0110912 Pulp Fiction (1994) 1994 9.0 490065 154 mins. Crime|
Thriller
tt0137523 Fight Club (1999) 1999 8.8 458173 139 mins. Drama|
Mystery|Thriller
tt0133093 The Matrix (1999) 1999 8.7 448114 136 mins. Action|
Adventure|Sci-Fi
tt1375666 Inception (2010) 2010 8.9 385149 148 mins. Action|
Adventure|Sci-Fi|Thriller
tt0109830 Forrest Gump (1994) 1994 8.7 368994 142 mins. Comedy|
Drama|Romance
tt0169547 American Beauty (1999) 1999 8.6 338332 122 mins. Drama
tt0499549 Avatar (2009) 2009 8.1 336855 162 mins. Action|
Adventure|Fantasy|Sci-Fi
tt0108052 Schindler's List (1993) 1993 8.9 325888 195 mins.
Biography|Drama|History|War
tt0080684 Star Wars: Episode V - The Empire Strikes Back
```

Here's a preview of the raw data we'll use -- it's a list of the 10,000 movies made since 1950 with the most IMDB user ratings.

```

%matplotlib inline

import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

names = ['imdbID', 'title', 'year', 'score', 'votes', 'runtime',
         'genres']
data = pd.read_csv('data/imdb_top_10000.txt', delimiter='\t',
                  names=names).dropna()

print "Number of rows: %i" % data.shape[0]
data.head() # print the first 5 rows

Number of rows: 9999
      imdbID      title  year  score
votes runtime genres
0  tt0111161  The Shawshank Redemption (1994)  1994    9.2
619479  142 mins.  Crime|Drama
1  tt0110912      Pulp Fiction (1994)  1994    9.0
490065  154 mins.  Crime|Thriller
...
```

The text file is tab-separated, and doesn't have any column headers. We set the appropriate keywords in `pd.read_csv()` to handle this and drop any empty values.

Notice the runtime is currently a string rather than an integer number. The genres column is not atomic. Finally the movie year is repeated in the year and title column.

```
dirty = '142 mins.'  
number, text = dirty.split(' ')  
clean = int(number)  
print number  
  
142  
  
clean_runtime = [float(r.split(' ')[0]) for r in data.runtime]  
data['runtime'] = clean_runtime
```

Let's fix the runtime column. The following snippet converts a string like '142 mins.' to the number 142. So we use it to replace the runtime column with the cleaned runtime.

Actually, lets do the rest in iPython. But before we continue let me talk a bit about matplotlib.

```
from matplotlib import figure
from matplotlib.backends.backend_agg import (FigureCanvasAgg as
FigureCanvas)

fig = figure.Figure()
canvas = FigureCanvas(fig)

canvas.print_figure('canvas.png', facecolor='lightgray')
```



29


matplotlib can be used in a procedural manner, as suggested by its name which was trying to copy matlab. But to get anything decent out of matplotlib you need to go object oriented.

So typically the way to use matplotlib is to define states in steps where we create a figure and add a canvas to it and then print it out. It's like a sandwich which makes it hard to pull out pieces in the middle to inspect or reuse them.

The object oriented way is to create objects and attach properties to different aspects of this object. So as shown above we will end up with just a blank light grey canvas.

```
fig = figure.Figure()
canvas = FigureCanvas(fig)
ax = fig.add_subplot(1,1,1)

canvas.print_figure('plot.png', facecolor='lightgray')
```



Now to add a plot we use the `add_subplot()` function to our figure. We always use the `add_subplot()` to add plots to our figure since it gives us control over where we put things on our canvas. `add_subplot()` takes the total number of rows, columns and the current plot's index.

```

import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

x1 = np.linspace(0.0, 5.0)
x2 = np.linspace(0.0, 2.0)

y1 = np.cos(2 * np.pi * x1) * np.exp(-x1)
y2 = np.cos(2 * np.pi * x2)

plt.subplot(2, 1, 1)
plt.plot(x1, y1, 'yo-')
plt.title('Oscillations')
plt.ylabel('Damped oscillation')

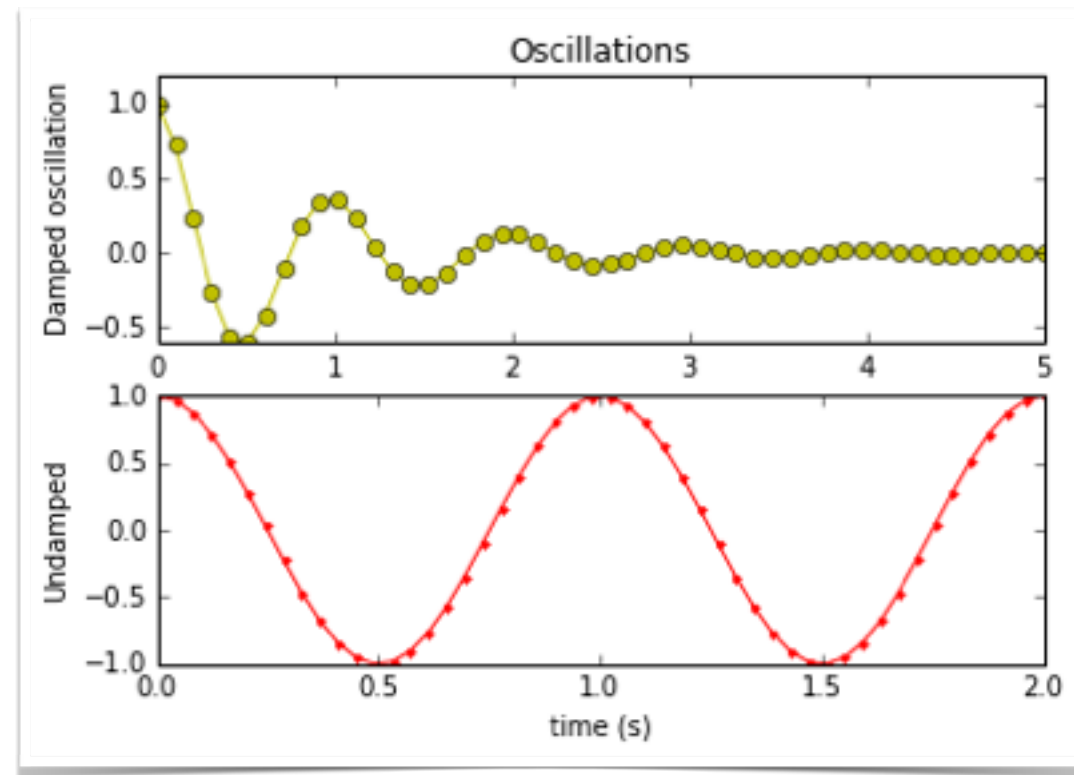
plt.subplot(2, 1, 2)
plt.plot(x2, y2, 'r.-')
plt.xlabel('time (s)')
plt.ylabel('Undamped')

plt.show()

```


31

pyplot provides a matlab like plotting environment. Here we use subplot to show two plots and configure the properties of the object like the plot label etc.



And this is what it looks like.





For further information check chapters 5 and 7 in this book.